

# Week 4 — Bayesian Generalization

Friday, May 22, 2026

Prof. Joseph Austerweil

# Agenda

|                                                        |      |
|--------------------------------------------------------|------|
| Welcome + Clusters walkthrough                         | 0:00 |
| The generalization problem                             | 0:10 |
| The Bayesian generalization framework + size principle | 0:18 |
| Rectangle game + number game                           | 0:40 |
| Break                                                  | 1:08 |
| Student presentation — Shohei                          | 1:15 |
| No Free Lunch                                          | 1:40 |
| Hierarchical Bayes + close                             | 1:50 |

# Assignment 1 — Clusters

# Clusters — what and when

## Assignment 1: Gaussians, Categories, and Clusters

- Due **Fri Jun 5, 2026, 8:00 PM** — worth 7.5%
- **Problem 1** — Gaussian-Gaussian conjugate model → *you can do this today* (it's Week 3 material)
- **Problem 2** — Gaussian mixture / categorization → Bayes' rule over two category Gaussians
- **Problem 3** — clustering → leans on this week + T3 Ch 5

# Clusters — which notebook

- `clusters.ipynb` is the canonical stencil — the GenJAX path
- `clusters_python.ipynb` and `clusters_nosoln.Rmd` — non-GenJAX paths, *same math, same credit*
- Matlab available on request
- “Open in Colab” links are live on the assignments page — one click, no local setup
- Read the **assignment PDF first** — it has the problem statements and all the math

Assignments page: [hml.chibatech.dev/assignments.html](http://hml.chibatech.dev/assignments.html)

# Where we are

|                                                        |      |
|--------------------------------------------------------|------|
| Welcome + Clusters walkthrough                         | 0:00 |
| <b>The generalization problem</b>                      | 0:10 |
| The Bayesian generalization framework + size principle | 0:18 |
| Rectangle game + number game                           | 0:40 |
| Break                                                  | 1:08 |
| Student presentation — Shohei                          | 1:15 |
| No Free Lunch                                          | 1:40 |
| Hierarchical Bayes + close                             | 1:50 |

# The generalization problem

# Chibany's lunches

Chibany has had three tonkatsu lunches this week. Each weighed about **500 g**.

Today a lunch arrives weighing **480 g**. Is it tonkatsu?

What about **700 g**? What about **350 g**?


# What just happened

**Generalization** — deciding when to extend a property from observed examples to a *novel* stimulus.

- No two stimuli are ever identical → you *must* generalize to act at all
- It is everywhere: word learning, categorization, object recognition, property induction, stereotypes
- It is the core problem of inductive inference

# Shepard's universal law

Shepard (1987) — across species and domains, the probability of generalization **decays exponentially** with distance in **psychological space** (*the perceived distance, after the mind represents the stimulus — not raw physical space*).



One law, one equation:

$$g(d) = e^{-d}$$

- $g$  — probability of generalization
- $d$  — distance in psychological space from the referent

The same curve holds across **species, senses, and stimulus types** — pitch, colour, line length, faces. Shepard called it *universal* for that reason.

But the law is **descriptive**: it says generalization decays exponentially, not *why*. Today's Bayesian framework will **derive** that exponential — it falls out of the posterior.

# Poll — Shepard's universal law

Shepard's universal law of generalization says generalization ...

- **A.** decays exponentially in psychological space
- **B.** decays exponentially in stimulus space
- **C.** grows exponentially in psychological space
- **D.** grows exponentially in stimulus space

# Poll — answer

**A. Decays exponentially in psychological space.**

“Stimulus space” is the trap — generalization isn’t governed by physical distance but by *perceived* distance. A model of generalization therefore needs a model of the **psychological space**. That is exactly what the Bayesian framework supplies next.

# Where we are

|                                                               |             |
|---------------------------------------------------------------|-------------|
| Welcome + Clusters walkthrough                                | 0:00        |
| The generalization problem                                    | 0:10        |
| <b>The Bayesian generalization framework + size principle</b> | <b>0:18</b> |
| Rectangle game + number game                                  | 0:40        |
| Break                                                         | 1:08        |
| Student presentation — Shohei                                 | 1:15        |
| No Free Lunch                                                 | 1:40        |
| Hierarchical Bayes + close                                    | 1:50        |

# The Bayesian generalization framework

# The idea — concepts as hypotheses

Instead of measuring distance directly, posit a **space of candidate concepts** and let Bayes do the generalizing.

- A concept = a **set of stimuli** that share the property
- Shepard's term: a **consequential subset** — the subset of things the property “applies to”
- A **feature** is a hypothesis too: “has stripes” picks out a set of stimuli — generalization is just asking which feature-sets the examples imply
- Generalization becomes: *which concepts are consistent with the examples, and do they contain the new stimulus?*

# Notation lock-in

- $h$  — a **hypothesis**: one candidate concept, i.e. a *set* of stimuli
- $\mathcal{H}$  — the **hypothesis space**: all candidate  $h$
- $X = \{x_1, \dots, x_n\}$  — the observed **examples** of the concept
- $y$  — a **novel stimulus** we must judge
- $C$  — the (unknown) true concept

# The three ingredients

**Prior**  $p(h)$  — domain knowledge: which concepts are *natural* before any data.

**Likelihood**  $p(X \mid h)$  — how probable the examples are if  $h$  is the true concept.

**Posterior**  $p(h \mid X)$  — belief in  $h$  after seeing the examples:

$$p(h \mid X) \propto p(X \mid h) p(h)$$

# The hypothesis space IS a prior

Choosing  $\mathcal{H}$  is already a **strong prior**.

Any concept not in  $\mathcal{H}$  has  $p(h) = 0$  — it can never be learned, no matter the data.

So: a learner's inductive bias lives in *which hypotheses it even considers*.

# Generalization = a posterior-weighted vote



Probability that the novel stimulus  $y$  has the property:

$$p(y \in C \mid X) = \sum_{h \in \mathcal{H}} \mathbf{1}[y \in h] p(h \mid X)$$

Every hypothesis votes. Its vote is **its posterior weight**, and it votes “yes” only if it *contains*  $y$ .

$\mathbf{1}[y \in h]$  — the indicator: 1 if  $y$  is in the set  $h$ , else 0.

# One datum → the posterior-weighted vote





**The hypotheses.** Observe one datum  $x$ . Every interval containing  $x$  is a live hypothesis  $h$ . Bar thickness =  $1/|h|$  (the strong-sampling likelihood); flat prior, so posterior  $\propto$  likelihood.

**The vote.** For each candidate  $y$ , sum the posterior of every  $h$  that contains it:

$$p(y \in C | x) = \sum_h \mathbf{1}[y \in h] p(h | x)$$

Next three slides: walk  $y$  outward —  $x$ ,  $x+1$ ,  $x+2$ .

# Vote for $y = x$




**The question.** Does the property generalize to  $y = x$ , the observed datum itself?

**Who votes.** Every hypothesis was built to contain  $x$ , so **all 7 of 7 vote “yes”**.

**The bar.** The vote sums *all* the posterior weight —  $\sum_h p(h | x) = 1$ . The **tallest bar the gradient can have**.

**The baseline.** This peak is what every other  $y$  is measured against — step away and the bar can only fall.


# Vote for $y = x + 1$



**The question.** Step one unit out: does the property generalize to  $y = x + 1$ ?


**Who drops.** The two smallest intervals no longer reach  $y$  — they grey out and vote “no”. **5 of 7 still vote “yes”.**

**The bar.** Fewer hypotheses contribute → a **shorter** bar than at  $x$ .



**Why it falls fast.** The dropouts are the **smallest** intervals — by the size principle the **heaviest-weight** votes. Losing those first is why the gradient decays steeply near  $x$ .


# Vote for $y = x + 2$



**The question.** Two steps out: does the property generalize to  $y = x + 2$ ?

**Who's left.** Only the 2 widest intervals still reach  $y$  — and those are the *least* likely (thinnest bars). **2 of 7 vote.**

**The bar.** Few votes, and only low-weight ones → the bar is short.



**The payoff.** Sweep  $y$  across every point and the bar heights trace an **approximately exponential decay** — Shepard's universal law, **derived** from the model, not assumed.

# The size principle

# Where did the examples come from?

The likelihood  $p(X \mid h)$  depends on *how you assume the examples were generated*. Two assumptions:

**Weak sampling** — examples generated some other way, then *labeled* by whether they fall in  $h$ .

**Strong sampling** — each example drawn *uniformly at random* from *within*  $h$ .

# The two likelihoods

## Weak sampling

$$p(X | h) = \begin{cases} 1 & \text{all } x_i \in h \\ 0 & \text{else} \end{cases}$$

Size-blind: a hypothesis either contains the data or it doesn't.

**Does not depend on  $|h|$ .**

$|h|$  — the size of hypothesis  $h$  (how many stimuli it contains).

$n$  — number of examples.

## Strong sampling

$$p(x | h) = \frac{1}{|h|} \Rightarrow p(X | h) = \left( \frac{1}{|h|} \right)^n$$

Each example drawn *uniformly from inside  $h$* .

**Smaller  $|h| \rightarrow$  higher likelihood, exponentially so in  $n$ .**

# The size principle

Under strong sampling: **smaller hypotheses get higher likelihood** — and *exponentially* more so as the number of examples  $n$  grows.

$$p(X \mid h) = \left( \frac{1}{|h|} \right)^n$$

So a small  $|h|$  wins — and **wins faster with more data**, because the exponent  $n$  magnifies any size advantage.

*The mechanism behind both games today.*

# Why — the suspicious coincidence

You see the examples {60, 80, 10, 30}.

- If the concept is “**multiples of 10**” — unremarkable, that’s just what its members look like.
- If the concept is “**even numbers**” — it’s a **suspicious coincidence** that not one of the four was 2, 4, 6, 8, ...

Strong sampling *penalizes* the big hypothesis for failing to predict the tight clustering you actually saw.

# Poll — strong sampling

What is strong sampling?

- **A.** Each stimulus is generated uniformly at random from the true hypothesis
- **B.** A stimulus has probability one given the true hypothesis
- **C.** Larger hypotheses are given smaller prior probability
- **D.** Smaller hypotheses are given smaller prior probability

# Poll — answer

**A. Uniformly at random from the true hypothesis.**

**B** describes *weak* sampling — membership gives likelihood 1, regardless of size. **C / D** describe a *prior* over hypotheses; the size principle is about the **likelihood**, not the prior. Strong sampling = “the examples were drawn *from inside* the concept” — and that is what makes size matter.

# Where we are

|                                                        |             |
|--------------------------------------------------------|-------------|
| Welcome + Clusters walkthrough                         | 0:00        |
| The generalization problem                             | 0:10        |
| The Bayesian generalization framework + size principle | 0:18        |
| <b>Rectangle game + number game</b>                    | <b>0:40</b> |
| Break                                                  | 1:08        |
| Student presentation — Shohei                          | 1:15        |
| No Free Lunch                                          | 1:40        |
| Hierarchical Bayes + close                             | 1:50        |

# Rectangle game — continuous concepts


# Gnarbles

A **gnarble** is a rectangle whose dimensions fall in some interval.

Tenenbaum's cover story: a **2-D concept** is an axis-aligned rectangle in a feature space — e.g. *healthy levels* of insulin × cholesterol.

You observe a few examples drawn from inside the rectangle. **Which other points are gnarbles?**

# Start in 1-D



Concept = an **interval**  $[\ell, u]$  on one dimension.

- $\mathcal{H}$  = *all* intervals
- Strong sampling  $\rightarrow$  likelihood  $\propto \left( \frac{1}{u - \ell} \right)^n$
- Hypothesis size = the **length**  $u - \ell$

Every interval that contains the data is a hypothesis; the **snug** ones get the most posterior weight.


# The generalization gradient

Plot  $p(y \in C \mid X)$  as  $y$  moves along the dimension:

- **High** inside the range of the examples
- **Decays** as  $y$  moves outside that range

The decay is *exponential* — Shepard's universal law, now **derived** rather than assumed.

# The 1-D gradient — built from votes




The same posterior-weighted vote as the number-line construction — now with **several examples  $X$** .

**Who votes.** Every interval that **contains all of  $X$**  is a live hypothesis; a snugger interval is thicker (more posterior).

**The gradient.** Sum the posterior of every interval that contains  $y$ :

$$p(y \in C \mid X) = \sum_h \mathbf{1}[y \in h] p(h \mid X)$$




Flat across the data, decaying outside — the gradient *is* the posterior-weighted vote, one  $y$  at a time.

# More examples → tighter generalization

- **One example** → broad, diffuse generalization (many interval sizes survive)
- **Many examples** → tight generalization, hugging the data range

Why: the size principle. With large  $n$ , big intervals lose likelihood *exponentially* fast — only the small, snug intervals keep posterior mass.

# Into 2-D — the rectangle game



Same machinery, one dimension up: a concept is an axis-aligned **rectangle**.

- Observe  $n$  dots inside the true rectangle
- Every rectangle that **encloses all**  $n$  is a hypothesis
- Smaller rectangle  $\rightarrow$  bigger likelihood (brighter / thicker)

$r$  — the range the data spans.


$d$  — how far a rectangle extends past that range.

# The rectangle experiment

Tenenbaum (1999) ran this as a behavioural experiment.

- On each trial, subjects saw  $n$  **dots** drawn from “an arbitrary rectangle of healthy insulin / cholesterol levels”
- They drew the rectangle they thought the dots came from
- $n$  varied from **2 to 50**; the data range  $r$  varied across trials
- The measure:  $d$  — how far past the data range  $r$  the drawn rectangle extends

# The result — $d$ vs. $r$ , by $n$




**Solid = human, dashed = model.** One colour per  $n$ .

**The human pattern.** Fewer examples  $\rightarrow$  generalize further ( $n = 2$  on top);  $d$  rises with  $r$  but **saturates**.

**The model — likelihood only.** The size principle with a *flat* (uninformative) prior. It captures the  $n$ -ordering...

...but the curves run **straight** — the model **over-extends**, badly for small  $n$  and large  $r$ . It misses the human saturation.

# One fix — an exponential prior



A flat prior lets the rectangle run **straight** — it **over-extends**. The fix: a prior that makes large rectangles less likely.


**The exponential distribution** — our first new distribution today. For a size  $s \geq 0$ :

$$p(s) = \lambda e^{-\lambda s}$$

- Always decreasing — **small  $s$  favoured**
- One parameter  $\lambda > 0$ ; the mean is  $1/\lambda$
- Larger  $\lambda \rightarrow$  faster decay  $\rightarrow$  stronger pull toward small rectangles

# With the prior — the fit

Human vs. model — expected-size (exponential) prior



**Same axes, same data** — now the model carries the exponential prior over rectangle size.

**The dashed curves bend.** The straight, over-extending lines collapse onto the human curves —  $d$  now **saturates** with  $r$ , just as people do.

Likelihood (size principle) **and** prior together — neither alone — give the fit.

Same lesson for every paper: lay the **model on top of the human data** and read off where it bends.

# The rectangle game in one line

The continuous concept learner is **just the framework equation** with  $\mathcal{H}$  = intervals / rectangles.

$$p(y \in C \mid X) = \sum_h \mathbf{1}[y \in h] p(h \mid X)$$

Nothing new — only the choice of  $\mathcal{H}$ .

# Number game — discrete concepts

# The number game

A simple task (Tenenbaum, 1999):

- I have a concept — a set of numbers between 1 and 100
- You see one or more “yes” examples
- You judge: is some other number a “yes”?

Watch what your own judgments do as examples come in.

# What people actually do

Human generalization judgments (Tenenbaum's  $N = 20$  subjects):

- Examples {60} → **diffuse** similarity: many numbers get moderate “yes”
- Examples {60, 80, 10, 30} → sharp “**multiples of 10**”
- Examples {60, 52, 57, 55} → sharp “**numbers near 60**”

One example → graded. Four examples → a crisp **rule**.

# Two things to explain

1. Generalization can look **similarity-based (graded)** or **rule-based (all-or-none)** — and people switch between them.
2. People learn a concept from **just a few examples**.

One model — the Bayesian framework — produces *both*, with no extra machinery.

# The discrete hypothesis space

$\mathcal{H}$  for the number game has two kinds of hypothesis:

**Mathematical properties** (~24) — even, odd, primes, squares, cubes, multiples of  $k$ , powers of  $k$ .

**Magnitude intervals** — “numbers in  $[a, b]$ ”: e.g. 10–20, 30–45.

The **prior**  $p(h)$  weights these families against each other.

# Size principle, by the numbers

Two candidate concepts for numbers in 1–100:

- **Multiples of 2** — 50 numbers  $\rightarrow p(x \mid h) = \frac{1}{50} = 2\%$  each
- **Multiples of 10** — 10 numbers  $\rightarrow p(x \mid h) = \frac{1}{10} = 10\%$  each

# One example: $x = 60$

$$p(60 \mid \text{mult-2}) = \frac{1}{50} \quad p(60 \mid \text{mult-10}) = \frac{1}{10}$$

Multiples of 10 is already **5×** more likely — but only 5×. With one example, many hypotheses stay in contention → **graded generalization**.

**Four examples: {10, 30, 60, 80}**

$$p(X \mid \text{mult-2}) = \left(\frac{1}{50}\right)^4 \approx 1.6 \times 10^{-7}$$

$$p(X \mid \text{mult-10}) = \left(\frac{1}{10}\right)^4 = 10^{-4}$$

Now multiples of 10 is **~625×** more likely. The 5× edge got raised to the 4th power → a crisp **rule**.

# From likelihood to posterior

Likelihoods aren't beliefs yet. Take a **two-hypothesis** model —  $\mathcal{H} = \{ \text{multiples of 10, even numbers} \}$ , both containing every example — with a flat prior, and turn the crank:


$$p(h \mid X) = \frac{p(X \mid h) p(h)}{\sum_{h'} p(X \mid h') p(h')}$$

Next: the posterior under **strong** vs. **weak** sampling, for  $X = \{60\}$   
then  $X = \{60, 80, 10, 30\}$ .

# Strong sampling — one example

$X = \{60\}$  · one example

Strong sampling



**The data:**  $X = \{60\}$  — a single example. Strong sampling:

$$p(h | X) \propto \left(\frac{1}{|h|}\right)^1 \cdot 0.5$$

- mult-10: likelihood  $\frac{1}{10}$  · even:  $\frac{1}{50}$
- Normalize the two → **0.83 vs. 0.17**


**Read it:** one example already **tilts** belief toward the smaller hypothesis — but only gently. A 5:1 likelihood ratio is not decisive, so plenty of posterior mass still sits on “even numbers”.

This is the **graded** regime — belief shifts, but no rule yet.

# Strong sampling — four examples

$X = \{60, 80, 10, 30\}$  · four examples

Strong sampling



$X = \{60, 80, 10, 30\}$ , strong sampling:

$$p(h | X) \propto \left(\frac{1}{|h|}\right)^4 \cdot 0.5$$


- The 5:1 ratio is now  $5^4 = 625:1$
- Normalize → **0.998 vs. 0.002**

Four “even” numbers, none of them odd-looking — that’s the **suspicious coincidence**. Strong sampling all but rules “even” out.

# Weak sampling — eliminate, but don't rank

both hypotheses survive  $\rightarrow$  posterior = prior (0.5 / 0.5)

Weak sampling



Weak-sampling likelihood is 1 **for any  $h$  that contains the data**, 0 otherwise — *no  $|h|$  dependence*.

**What it can do:** a datum *outside*  $h$  kills it — likelihood 0. Weak sampling **does** move the posterior, by **ruling hypotheses out**.

**What it can't do:** among the hypotheses that *all still contain the data*, it has no preference — every survivor keeps likelihood 1, so the posterior over them stays at the **prior**.

Here both hypotheses contain every example — nothing is ruled out:

$$p(h | X) = \frac{1 \cdot 0.5}{1 \cdot 0.5 + 1 \cdot 0.5} = 0.5$$

So weak sampling can't see the suspicious coincidence — it can't tell “multiples of 10” from “even numbers”.

# The number game in one line

Same equation as the rectangle game —  $\mathcal{H}$  is now **discrete**.

$$p(y \in C \mid X) = \sum_h \mathbf{1}[y \in h] p(h \mid X)$$

Graded vs. rule-like generalization both **fall out of the posterior** — no extra mechanism. The size principle's exponent does the switching.

# Where we are

|                                                        |             |
|--------------------------------------------------------|-------------|
| Welcome + Clusters walkthrough                         | 0:00        |
| The generalization problem                             | 0:10        |
| The Bayesian generalization framework + size principle | 0:18        |
| Rectangle game + number game                           | 0:40        |
| <b>Break</b>                                           | <b>1:08</b> |
| Student presentation — Shohei                          | 1:15        |
| No Free Lunch                                          | 1:40        |
| Hierarchical Bayes + close                             | 1:50        |

**Break**

# Bridge — Shohei's paper

Up next: Shohei presents **Tenenbaum & Xu (2000), *Word learning as Bayesian inference***.

A child hears “*this is a dax*” pointing at three Dalmatians. Is a poodle a dax? A cat?

This is the **number game**, with  $\mathcal{H}$  = candidate word meanings (subordinate / basic-level / superordinate). The **size principle** explains why three subordinate examples → a subordinate meaning.

*Watch for the size principle doing the work.*

# Student presentation — Shohei

# Where we are

|                                                        |             |
|--------------------------------------------------------|-------------|
| Welcome + Clusters walkthrough                         | 0:00        |
| The generalization problem                             | 0:10        |
| The Bayesian generalization framework + size principle | 0:18        |
| Rectangle game + number game                           | 0:40        |
| Break                                                  | 1:08        |
| Student presentation — Shohei                          | 1:15        |
| <b>No Free Lunch</b>                                   | <b>1:40</b> |
| Hierarchical Bayes + close                             | 1:50        |

# No Free Lunch

# The No Free Lunch theorem

Wolpert (1996) — averaged over **all possible worlds**, no learning algorithm beats any other. Here is *why*, concretely.

**The task.** You see the bits 0, 1 and must predict  $x_3$ .

**They pair off.** Every world where your rule scores a point has a mirror world — identical data, flipped continuation — where it loses one.

**Your rule.** Say it predicts  $x_3 = 0$ . In the world 0, 1, 0 it is **right**.

**Sum the pair.** Right + wrong = 1 hit out of 2. Average over *all* worlds: every algorithm, every rule, scores **exactly** 1/2.

**The mirror world.** But the world 0, 1, 1 is *just as possible* — same data 0, 1, opposite answer. There your rule is **wrong**.

No rule can win the average, because the data 0, 1 says **nothing** about  $x_3$  until you assume some worlds are more likely than others.

# What NFL means for us

A learner only works because the distribution over worlds is **constrained** — i.e. because it has a **non-flat prior**.

- Generalization is *impossible* without inductive bias
- Recall: the hypothesis space  $\mathcal{H}$  and the prior  $p(h)$  from Block 3
- They are not bookkeeping — they are *the entire reason* learning is possible

# Poll — No Free Lunch

What is the No Free Lunch theorem (for prediction)?

- **A.** When all hypotheses are possible, there's nothing you can learn to predict
- **B.** Learning one hypothesis hurts learning other hypotheses
- **C.** If someone gives you lunch for free, they'll expect something back
- **D.** Generalizing to new stimuli can hurt a learner

# Poll — answer

**A. When all hypotheses are possible, there's nothing you can learn to predict.**

A flat prior over every possible world means data carries no leverage — every continuation stays 50/50. To learn, you must *commit* to some worlds being more likely than others. **That commitment is your prior — and it is why your prior matters.**

# Where we are

|                                                        |             |
|--------------------------------------------------------|-------------|
| Welcome + Clusters walkthrough                         | 0:00        |
| The generalization problem                             | 0:10        |
| The Bayesian generalization framework + size principle | 0:18        |
| Rectangle game + number game                           | 0:40        |
| Break                                                  | 1:08        |
| Student presentation — Shohei                          | 1:15        |
| No Free Lunch                                          | 1:40        |
| <b>Hierarchical Bayes + close</b>                      | <b>1:50</b> |

# Hierarchical Bayes

# Back to Chibany's class

Every student in Chibany's class has their **own** tonkatsu-vs-hamburger rate  $\theta_i$ .

The rates aren't identical — but they aren't unrelated either. They're all Chibany's customers.

Should learning Aoi's rate tell us anything about Ben's?

# Priors over priors

Put a prior on the **parameters of the prior** — and treat  $(a, b)$  themselves as unknown:

$$\theta_i \sim \text{Beta}(a, b)$$

The two-level model:  $(a, b) \longrightarrow \theta_i \longrightarrow$  observed bentos.

$(a, b)$  is the *shared* structure; each  $\theta_i$  is a student.

# Why it matters

A hierarchical model lets a learner **learn the prior from data** — exactly the inductive bias No Free Lunch said you cannot do without.

- $(a, b)$  is learned from *all* the students together
- It is how “overhypotheses” get acquired — the shape bias, object vs. substance kinds (Kemp, Perfors & Tenenbaum, 2007)

# Three ways to pool the data

Six students, six bento records. How do you estimate their rates?

| Approach         | Model                                 | Problem                                      |
|------------------|---------------------------------------|----------------------------------------------|
| Complete pooling | one shared $\theta$ for everyone      | ignores that students differ                 |
| No pooling       | a separate $\theta_i$ , all unrelated | ignores that they're all Chibany's customers |
| Hierarchical     | $\theta_i \sim \text{Beta}(a, b)$     | <i>the middle path</i> — borrows strength    |

Hierarchical Bayes = **partial pooling**.

# The two-level model

Built up, top to bottom:

$$(a, b) \sim \text{prior}$$

$$\theta_i \mid a, b \sim \text{Beta}(a, b)$$

$$k_i \mid \theta_i \sim \text{Binomial}(n_i, \theta_i)$$

$k_i$  — tonkatsu count for student  $i$ .  $n_i$  — that student's total bentos.

# Inference — no closed form

We want the posterior over *everything*:

$$p(a, b, \{\theta_i\} \mid \text{data})$$

Unlike the Beta-Binomial of Week 3, this has **no clean closed form**.

This is where sampling comes in — and where **GenJAX** earns its place (the hierarchical `bento_day()` exercise builds exactly this model).

# Shrinkage — borrowing strength

For each student, compare:

- their **raw** tonkatsu fraction  $k_i/n_i$
- their **posterior mean**  $\hat{\theta}_i$

The posterior means are pulled **toward the group mean** — *a lot* when a student has little data, barely at all when they have a lot.

This is hierarchical Bayes **automatically borrowing strength** across students.

# Overhypotheses — learning the bias

With the hierarchy in hand, **overhypotheses** become precise.

The shape bias; the object-vs-substance distinction — these are *second-level* hypotheses, learned as a distribution over *kinds of concept* (Kemp, Perfors & Tenenbaum, 2007).

No Free Lunch said a learner needs inductive bias. The hierarchy is where a learner **acquires** it — instead of being born with it.

**Close**

# Before next week

Read **T3 Ch 5 — Mixture models** before Week 5.

It formalizes the same partial-pooling idea you just saw, and feeds directly into **Problem 3** of the Clusters assignment.

Clusters is due **Fri Jun 5, 8:00 PM**.

# Next week — Week 5

**Bayes nets + causal Bayes nets.**

From “concepts as sets” to **structured probabilistic models** — graphs that encode how variables depend on each other, and what happens when you *intervene*.

Check the readings page for the Week 5 paper + presenter.